

Image segmentation and synthesis

Quiz, 7 questions

1
point

1.

Which of the following is an operation, not a task in computer vision?

- ☐ Object detection
 - ☒ Perspective projection
 - ☐ Instance segmentation
 - ☒ Max-pooling
 - ☒ Gradient computation
 - ☒ Image convolution
-

1
point

2.

For a 3-class semantic segmentation problem, how many numbers must an algorithm output for a 640x480 image?

921600

1
point

3.

Why is SLIC algorithm better suited to the image oversegmentation task than k-means method?

- ☒ It is more computationally efficient because segment sizes are bounded, limiting the number of pixels examined at each iteration
- ☐ Unlike k-means, SLIC is a supervised learning method and thus can use labels to improve segmentation
- ☐ It utilizes a more robust distance metric, rather than simple Euclidean distance used in k-means method
- ☒



It limits distance between pixels by a certain threshold, utilizing the notion of hard spatial neighbourhood

Image segmentation and synthesis

Quiz, 7 questions

1
point

4.

What is the goal of the unpooling operation?

- ☐ To undo channel concatenation by decreasing the number of convolutional feature maps
 - ☐ To undo convolution by applying the transposed convolution
 - ☐ To help backpropagate errors by introducing sparse convolutions
 - ☒ To undo pooling by outputting an image with larger resolution (i.e., pixels in spatial directions)
-

1
point

5.

In unpooling, how do we approximate the inverse of the non-invertible max-pooling operation?

- ☒ We output maximal values at their respective indexes (called max location switches) and place zeroes elsewhere
 - ☐ We do bilinear interpolation to compute the output
 - ☐ We use 'bed of nails': output the maximal values in the top left corner and zeros elsewhere
-

1
point

6.

What is a Gram matrix in linear algebra?

- ☒ A matrix produced by computing dot product between two sets of vectors
 - ☐ A matrix of feature activations in a CNN
 - ☐ A confusion matrix of CNN
 - ☐ A positive-semidefinite matrix used to generate random numbers from a Gaussian distribution
-

1

point

Image segmentation and synthesis

Quiz, 7 questions

What makes a good generator for a GAN model?

- ☐ It produces nicely looking images
- ☐ It achieves superior performance in generating Gaussian mixtures
- ☒ It produces data that is hard to distinguish from real

☐ I, **Jiadaï Zhao**, understand that submitting work that isn't my own may result in permanent failure of this course or deactivation of my Coursera account.

[Learn more about Coursera's Honor Code](#)

Submit Quiz

